

1. When conducting the work of Lab 11, we conducted the test that uses the Central Limit Theorem even though the sample size was “small” (i.e., $n < 30$). It turns out, that how “far off” the t -test is can be computed using a first-order Edgeworth approximation for the error. Below, we will do this for the the further observations.

(a) Boos and Hughes-Oliver (2000) note that

$$P(T \leq t) \approx F_Z(t) + \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6}}_{\text{error}} f_Z(t),$$

where $f_Z(\cdot)$ and $F_Z(\cdot)$ are the Gaussian PDF and CDF and skew is the skewness of the data. What is the potential error in the computation of the p -value when testing $H_0 : \mu_X = 0; H_a : \mu_X < 0$ using the zebra finch further data?

```
## 1a
zebra.fitches.dat <- read_csv("zebrafinches.csv")

further.vec <- zebra.fitches.dat$further
skew <- skewness(further.vec)
n <- length(further.vec)

further.test <- t.test(x=further.vec, mu = 0, alternative = "less")
t.far <- further.test$statistic #finding the t value
fz <- dnorm(t.far)
Fz <- pnorm(t.far)
(error = (skew / sqrt(n)) * ((2 * (t.far^2) + 1) / 6) * fz)

##          t
## -1.226006e-13

prob <- error + Fz
```

- (b) Compute the error for t statistics from -10 to 10 and plot a line that shows the error across t . Continue to use the skewness and the sample size for the zebra finch further data.

```
t.vecs <- seq(-10,10,0.1)
results <- tibble(
  t = t.vecs,
  fz = dnorm(t.vecs),
  error = (skew / sqrt(n)) * ((2 * (t.vecs^2) + 1) / 6) * dnorm(t.vecs)
)

error.plot <- ggplot(results) +
  theme_bw() +
  geom_line(aes(x = t, y = error)) +
  labs(
    x = "T values",
    title = "Edgeworth Approximation Error"
  )
```

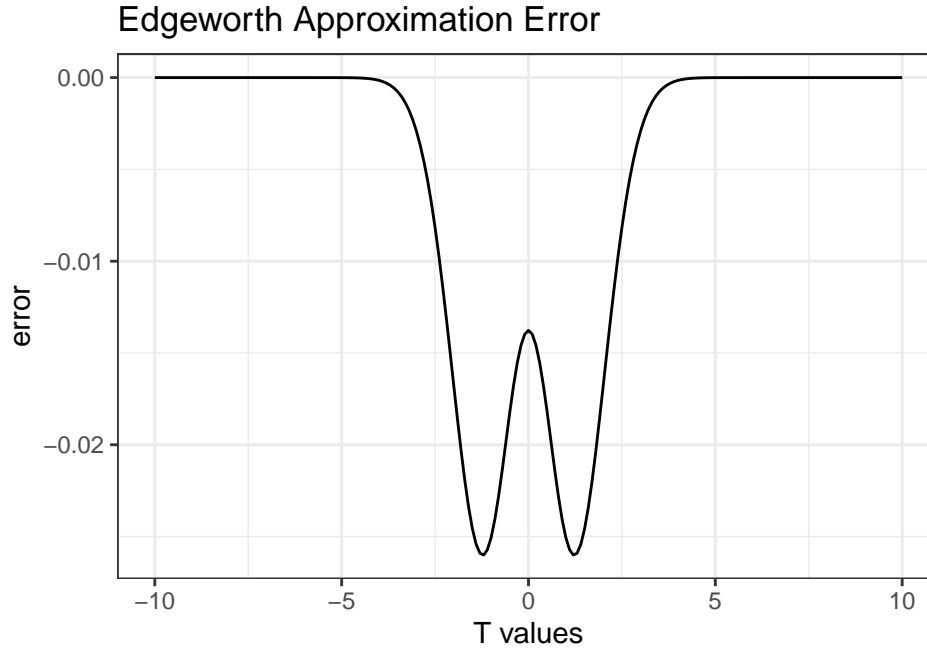


Figure 1: Edgeworth Error across span of t values

- (c) Suppose we wanted to have a tail probability within 10% of the desired $\alpha = 0.05$. Recall we did a left-tailed test using the further data. How large of a sample size would we need? That is, we need to solve the error formula equal to 10% of the desired left-tail probability:

$$0.10\alpha \stackrel{\text{set}}{=} \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6} f_Z(t)}_{\text{error}},$$

which yields

$$n = \left(\frac{\text{skew}}{6(0.10\alpha)} (2t^2 + 1) f_Z(t) \right)^2.$$

Answer: We would need a sample size of at least 521 for a left-tailed test.

```
alpha = 0.05
t.val <- qnorm(alpha)
pdf <- dnorm(t.val)
(target.size = ((skew / (6*(0.1*alpha))) * (2 * (t.val^2) + 1) * pdf)^2)
## [1] 520.8876
```

2. Complete the following steps to revisit the analyses from lab 11 using the bootstrap procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform resampling to approximate the sampling distribution of the T statistic:

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}},$$

where \bar{x}_r is the mean computed on the r^{th} resample and s is the sample standard deviation from the original samples. At the end, create an object called `resamples.null.closer`, for example, and store the resamples shifted to ensure they are consistent with the null hypotheses at the average (i.e., here ensure the shifted resamples are 0 on average, corresponding to $t = 0$, for each

case).

Answer: All of the shifted means of the t values are so small that they are practically 0.

```
#####
# Part 2: Boot strapping #
#####
R <- 10000
mu0 = 0
fur.sd <- sd(zebra.finches.dat$farther)
close.sd <- sd(zebra.finches.dat$closer)
diff.sd <- sd(zebra.finches.dat$diff)

resample <- tibble(close.t=numeric(R),
  far.t=numeric(R),
  diff.t=numeric(R),
  close.mean =numeric(R),
  far.mean =numeric(R),
  diff.mean =numeric(R)
)
resamples.shifted <- tibble()

for(i in 1:R){
  close.sample <- sample(x=zebra.finches.dat$closer,
    size=n,
    replace=T)
  far.sample <- sample(x=zebra.finches.dat$farther,
    size=n,
    replace=T)
  diff.sample <- sample(x=zebra.finches.dat$diff,
    size=n,
    replace=T)

  resample$close.t[i] = (mean(close.sample)-mu0)/(close.sd/sqrt(n))
  resample$far.t[i] = (mean(far.sample)-mu0)/(fur.sd/sqrt(n))
  resample$diff.t[i] = (mean(diff.sample)-mu0)/(diff.sd/sqrt(n))
  resample$close.mean[i] = mean(close.sample)
  resample$far.mean[i] = mean(far.sample)
  resample$diff.mean[i] = mean(diff.sample)
}

## shifting the means
close.change <- mean(resample$close.t)
far.change <- mean(resample$far.t)
diff.change <- mean(resample$diff.t)

resamples.shifted <- resample |>
  mutate(close.shifted = close.t - close.change) |>
  mutate(far.shifted = far.t - far.change) |>
  mutate(diff.shifted = diff.t - diff.change) |>
  select(c(close.shifted, far.shifted, diff.shifted ))
(mean(resamples.shifted$close.shifted))

## [1] 1.171951e-16

(mean(resamples.shifted$far.shifted))

## [1] 1.382103e-16

(mean(resamples.shifted$diff.shifted))

## [1] 3.648193e-16
```

- (b) Compute the bootstrap p -value for each test using the shifted resamples. How do these compare to the t -test p -values?

Answer: Table 1 below shows that the comparison of the p -values. While t -tests did produce nonzero p -values, they were so small that they are essentially 0, which is what the bootstrap test shows for each case.

```
# Bootstrap P-Value
boot.p.closer <- mean(resamples.shifted$close.shifted >= close.change)
boot.p.far <- mean(resamples.shifted$far.shifted <= far.change)
low = -diff.change
high = diff.change
p.low <- mean(resamples.shifted$diff.shifted <= low)
p.high <- mean(resamples.shifted$diff.shifted >= high)
boot.p.diff <- p.low +p.high
```

```
# T-Test P-Value
test.close <- t.test(zebra.finch.dat$closer,
                    mu = 0,
                    alternative = "greater")
close.ttest.p <- test.close$p.value

test.far <- t.test(zebra.finch.dat$further,
                  mu = 0,
                  alternative = "less")
far.ttest.p <- test.far$p.value

test.diff <- t.test(zebra.finch.dat$diff,
                   mu = 0,
                   alternative = "two.sided")
diff.ttest.p <- test.diff$p.value

comparisons.p <- tibble(" " = c("close", "far", "diff"), Bootstrapped =
                        c(boot.p.close, boot.p.far, boot.p.diff),
                        t.test = c(close.ttest.p, far.ttest.p, diff.ttest.p))
```

Data	Bootstrapped	T test
closer	0.00	0.00
further	0.00	0.00
difference	0.00	0.00

Table 1: Comparison of p-values between bootstrapping and T test

- (c) What is the 5th percentile of the shifted resamples under the null hypothesis?

Answer: Table 2 below shows that the comparison of the 5th percentile of the shifted resamples: closer = -1.54, further = -1.70, and difference = -1.61. They approximate the T-test, which is -1.71.

```
## Part c
close.5th <- quantile(resamples.shifted$close.shifted, 0.05)
far.5th <- quantile(resamples.shifted$far.shifted, 0.05)
diff.5th <- quantile(resamples.shifted$diff.shifted, 0.05)

close.percentile.t <- qt(0.05, df = n-1)
far.percentile.t <- qt(0.05, df = n-1)
diff.percentile.t <- qt(0.05, df = n-1)

comparisons.percentiles <- tibble(" " = c("close", "far", "diff"), sampled =
                                c(close.5th, far.5th, diff.5th),
                                t.val = c(close.percentile.t, far.percentile.t,
                                           diff.percentile.t))
```

Data	Bootstrapped	T-test
Closer	-1.54	-1.71
Further	-1.70	-1.71
Difference	-1.61	-1.71

Table 2: Comparison of the 5th percentile between bootstrapping and T test

Note this value approximates $t_{0.05, n-1}$. Compare these values in each case.

- (d) Compute the bootstrap confidence intervals using the resamples. How do these compare to the t -test confidence intervals?

Answer: Table 3 below shows that the comparison of the confidence intervals of the shifted resamples with those of the T-tests. Respectively for the closer, further, and difference cases, bootstrapping produced 95% confidence intervals of (0.12, 0.18), (-0.26, -0.17), and (0.27, 0.42) while the T-tests produced CIs of (0.12, 0.20), (-0.26, -0.15), and (0.27, 0.45). They are pretty similar to each other for each case

```
## Part d
lower.close <- quantile(resample$close.mean, 0.025)
upper.close <- quantile(resample$close.mean, 0.925)
Bootstrap.CI.close <- c(lower.close, upper.close)

lower.far <- quantile(resample$far.mean, 0.025)
upper.far <- quantile(resample$far.mean, 0.925)
Bootstrap.CI.far <- c(lower.far, upper.far)

lower.diff <- quantile(resample$diff.mean, 0.025)
upper.diff <- quantile(resample$diff.mean, 0.925)
Bootstrap.CI.diff <- c(lower.diff, upper.diff)

# t tests confidence intervals
test.close <- t.test(zebra.finches.dat$closer,
  mu = 0, conf.level = 0.95,
  alternative = "two.sided")
close.CI <- test.close$conf.int

test.far <- t.test(zebra.finches.dat$further,
  mu = 0, conf.level = 0.95,
  alternative = "two.sided")
far.CI <- test.far$conf.int

test.diff <- t.test(zebra.finches.dat$diff,
  mu = 0, conf.level = 0.95,
  alternative = "two.sided")
diff.CI <- test.diff$conf.int

comparisons.CI <- tibble(
  "Condition" = c("close", "far", "diff"),
  "Bootstrapped Lower" = c(lower.close, lower.far, lower.diff),
  "Bootstrapped Upper" = c(upper.close, upper.far, upper.diff),
  "t-test Lower" = c(close.CI[1], far.CI[1], diff.CI[1]),
  "t-test Upper" = c(close.CI[2], far.CI[2], diff.CI[2]))
```

Data	Bootstrapped Lower	Bootstrapped Upper	t-test Lower	t-test Upper
close	0.12	0.18	0.12	0.20
far	-0.26	-0.17	-0.26	-0.15
diff	0.27	0.42	0.27	0.45

Table 3: Comparison of the Confidence Interval between bootstrapping and T test, separated into lower and upper bounds

3. Complete the following steps to revisit the analyses from lab 11 using the randomization procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform the randomization procedure

```
mu0 <- 0
R <- 10000
randomized <- tibble(close.means = numeric(R), far.means = numeric(R),
  diff.means = numeric(R))

shifted.x.close <- zebra.finches.dat$closer - mu0
shifted.x.far <- zebra.finches.dat$further - mu0
shifted.x.diff <- zebra.finches.dat$diff - mu0

for(i in 1:R){
  curr.rand1 <- shifted.x.close *
    sample(x = c(-1, 1),
      size = length(shifted.x.close),
      replace = T)
  curr.rand2 <- shifted.x.far *
    sample(x = c(-1, 1),
      size = length(shifted.x.far),
      replace = T)
  curr.rand3 <- shifted.x.diff *
    sample(x = c(-1, 1),
      size = length(shifted.x.diff),
      replace = T)

  randomized$close.means[i] <- mean(curr.rand1)
  randomized$far.means[i] <- mean(curr.rand2)
```

```

    randomized$diff.means[i] <- mean(curr.rand3)
  }

randomized <- randomized |>
  mutate(close.means = close.means + mu0,
         far.means = far.means + mu0,
         diff.means = diff.means + mu0 ) # shifting data back

```

- (b) Compute the randomization test p -value for each test.

Answer: The p value is 0 for all three cases, so the close, further, and difference data are in the tailed regions of their respective null distributions. The closer data is in its right-tailed zone, further in its left-tailed zone, and difference in both tailed zones, so they all reject the null hypothesis.

```

#####
## 3b
## found observed data
observed.close.mean <- mean(zebra.finches.dat$closer -mu0)
observed.far.mean <- mean(zebra.finches.dat$further -mu0)
observed.diff.mean <- mean(zebra.finches.dat$diff -mu0)

## p-values for each value type
rand.p.close <- mean(randomized$close.means >= observed.close.mean)
rand.p.far <- mean(randomized$far.means <= observed.far.mean)

(delta <- abs(mean(zebra.finches.dat$diff) - mu0))

## [1] 0.3589475

(low <- mu0 - delta) # mirror

## [1] -0.3589475

(high<- mu0 + delta) # xbar

## [1] 0.3589475

rand.p.diff <- mean(randomized$diff.means <= low) +
  mean(randomized$diff.means >= high)

c(rand.p.close,rand.p.far,rand.p.diff)

## [1] 0 0 0

```

- (c) Compute the randomization confidence interval by iterating over values of μ_0 .

Hint: You can “search” for the lower bound from Q_1 and subtracting by 0.0001, and the upper bound using Q_3 and increasing by 0.0001. You will continue until you find the first value for which the two-sided p -value is greater than or equal to 0.05.

Answer: We have to numerically compute the confidence interval of a randomization test, so we slowly decreased the mean until the p -value hit 0.05 to find the lower bound and slowly increased the mean until the same condition to find the upper bound for each case. Our final result showed a confidence interval for the closer data as (0.1189, 0.1941), for the further data (-0.2627, -0.1427), and for the difference data (0.2689, 0.4489).

```

#####
## 3c
## creating randomized Confidence Intervals

R <- 1000
mu0.iterate <- 0.0001
starting.point <- mean(zebra.finches.dat$closer)
mu.lower.close <- starting.point

repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- zebra.finches.dat$closer - mu.lower.close
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

```

```

    rand$xbars[i] <- mean(curr.rand)
  }

  rand <- rand |>
  mutate(xbars = xbars + mu.lower.close) # shifting back

  # p-value
  delta <- abs(mean(zebra.finches.dat$closer) - mu.lower.close)
  low <- mu.lower.close - delta # mirror
  high <- mu.lower.close + delta # xbar
  p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high)

  if(p.val < 0.05){
    break
  }else{
    mu.lower.close <- mu.lower.close - mu0.iterate
  }
}

mu.upper.close <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- zebra.finches.dat$closer - mu.upper.close
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }

  rand <- rand |>
  mutate(xbars = xbars + mu.upper.close) # shifting back

  # p-value
  delta <- abs(mean(zebra.finches.dat$closer) - mu.upper.close)
  (low <- mu.upper.close - delta) # mirror
  (high <- mu.upper.close + delta) # xbar
  (p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high))

  if(p.val < 0.05){
    break
  }else{
    mu.upper.close <- mu.upper.close + mu0.iterate
  }
}

## CI for further
starting.point <- mean(zebra.finches.dat$further)
mu.lower.far <- starting.point

repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- zebra.finches.dat$further - mu.lower.far
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }

  rand <- rand |>
  mutate(xbars = xbars + mu.lower.far) # shifting back

  # p-value
  delta <- abs(mean(zebra.finches.dat$further) - mu.lower.far)

```

```

low <- mu.lower.far - delta # mirror
high<- mu.lower.far + delta # xbar
p.val <- mean(rand$xbars <= low) +
  mean(rand$xbars >= high)

if(p.val < 0.05){
  break
}else{
  mu.lower.far <- mu.lower.far - mu0.iterate
}
}

mu.upper.far <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- zebra.finch.dat$farther - mu.upper.far
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }

  rand <- rand |>
    mutate(xbars = xbars + mu.upper.far) # shifting back

  # p-value
  delta <- abs(mean(zebra.finch.dat$farther) - mu.upper.far)
  low <- mu.upper.far - delta # mirror
  high<- mu.upper.far + delta # xbar
  p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high)

  if(p.val < 0.05){
    break
  }else{
    mu.upper.far <- mu.upper.far + mu0.iterate
  }
}

## CI for diff
starting.point <- mean(zebra.finch.dat$diff)
mu.lower.diff <- starting.point

repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- zebra.finch.dat$diff - mu.lower.diff
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
            size = length(x.shift),
            replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }

  rand <- rand |>
    mutate(xbars = xbars + mu.lower.diff) # shifting back

  # p-value
  delta <- abs(mean(zebra.finch.dat$diff) - mu.lower.diff)
  low <- mu.lower.diff - delta # mirror
  high<- mu.lower.diff + delta # xbar
  p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high)

  if(p.val < 0.05){
    break
  }
}

```



```

    }else{
      mu.lower.diff <- mu.lower.diff - mu0.iterate
    }
  }

mu.upper.diff <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- zebra.finch.dat$diff - mu.upper.diff
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
             size = length(x.shift),
             replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }

  rand <- rand |>
    mutate(xbars = xbars + mu.upper.diff) # shifting back

  # p-value
  delta <- abs(mean(zebra.finch.dat$diff) - mu.upper.diff)
  low <- mu.upper.diff - delta # mirror
  high <- mu.upper.diff + delta # xbar
  p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high)

  if(p.val < 0.05){
    break
  }else{
    mu.upper.diff <- mu.upper.diff + mu0.iterate
  }
}

rand.CI <- tibble(" " = c("close", "far", "diff"),
                  lower.limit = c(mu.lower.close, mu.lower.far, mu.lower.diff),
                  upper.limit = c(mu.upper.close, mu.upper.far, mu.upper.diff))

```

4. **Optional Challenge:** In this lab, you performed resampling to approximate the sampling distribution of the T statistic using

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}}.$$

I'm curious whether it is better/worse/similar if we computed the statistics using the sample standard deviation of the resamples (s_r), instead of the original sample (s)

$$T = \frac{\bar{x}_r - 0}{s_r/\sqrt{n}}.$$

- Perform a simulation study to evaluate the Type I error for conducting this hypothesis test both ways.
- Using the same test case(s) as part (a), compute bootstrap confidence intervals and assess their coverage – how often do we ‘capture’ the parameter of interest?

References

Boos, D. D. and Hughes-Oliver, J. M. (2000). How large does n have to be for z and t intervals? *The American Statistician*, 54(2):121–128.